

Week 06: Graph Data Structures

Graph Definitions

Graphs

2/67

Many applications require

- a collection of *items* (i.e. a set)
- *relationships*/connections between items

Examples:

- maps: items are cities, connections are roads
- web: items are pages, connections are hyperlinks

Collection types you're familiar with

- lists ... linear sequence of items (week 3, COMP9021)
- trees ... branched hierarchy of items (COMP9021)

Graphs are more general ... allow arbitrary connections


... Graphs

3/67

A graph $G = (V, E)$

- V is a set of *vertices*
- E is a set of *edges* (subset of $V \times V$)

Example:



... Graphs

4/67

A real example: Australian road distances

Distance	Adelaide	Brisbane	Canberra	Darwin	Melbourne	Perth	Sydney
Adelaide	-	2055	1390	3051	732	2716	1605


Brisbane	2055	-	1291	3429	1671	4771	982
Canberra	1390	1291	-	4441	658	4106	309
Darwin	3051	3429	4441	-	3783	4049	4411
Melbourne	732	1671	658	3783	-	3448	873
Perth	2716	4771	4106	4049	3448	-	3972
Sydney	1605	982	309	4411	873	3972	-

Notes: vertices are cities, edges are distance between cities, symmetric

... Graphs

5/67

Alternative representation of above:



... Graphs

6/67

Questions we might ask about a graph:

- is there a way to get from item A to item B?
- what is the best way to get from A to B?
- which items are connected?

Graph algorithms are generally more complex than tree/list ones:

- no implicit order of items
- graphs may contain cycles
- concrete erpresentation is less obvious
- algorithm complexity depends on connection complexity

Properties of Graphs

7/67

Terminology: $|V|$ and $|E|$ (cardinality) normally written just as V and E .

A graph with V vertices has at most $V(V-1)/2$ edges.

The ratio $E:V$ can vary considerably.

- if E is closer to V^2 , the graph is *dense*
- if E is closer to V , the graph is *sparse*
 - Example: web pages and hyperlinks

Knowing whether a graph is sparse or dense is important

- may affect choice of data structures to represent graph
- may affect choice of algorithms to process graph

Exercise #1: Number of Edges

8/67

The edges in a graph represent pairs of connected vertices. A graph with V has V^2 such pairs.

Consider $V = \{1,2,3,4,5\}$ with all possible pairs:

$$E = \{ (1,1), (1,2), (1,3), (1,4), (1,5), (2,1), (2,2), \dots, (4,5), (5,5) \}$$

Why do we say that the maximum #edges is $V(V-1)/2$?

... because

- (v,w) and (w,v) denote the same edge (in an undirected graph)
- we do not consider loops (v,v)

Graph Terminology

10/67

For an edge e that connects vertices v and w

- v and w are *adjacent* (neighbours)
- e is *incident* on both v and w

Degree of a vertex v

- number of edges incident on e

Synonyms:

- vertex = node, edge = arc = link (Note: some people use arc for *directed* edges)

... Graph Terminology

11/67

Path: a sequence of vertices where


- each vertex has an edge to its predecessor

Cycle: a path where

- last vertex in path is same as first vertex in path

Length of path or cycle:

- #edges



... Graph Terminology

12/67

Connected graph

- there is a *path* from each vertex to every other vertex
- if a graph is not connected, it has ≥ 2 *connected components*

Complete graph K_V

- there is an *edge* from each vertex to every other vertex
- in a complete graph, $E = V(V-1)/2$



... Graph Terminology

13/67

Tree: connected (sub)graph with no cycles

Spanning tree: tree containing all vertices

Clique: complete subgraph

Consider the following single graph:



This graph has 25 vertices, 32 edges, and 4 connected components

Note: The entire graph has no spanning tree; what is shown in green is a spanning tree of the third connected component

... Graph Terminology

14/67

A *spanning tree* of connected graph $G = (V, E)$

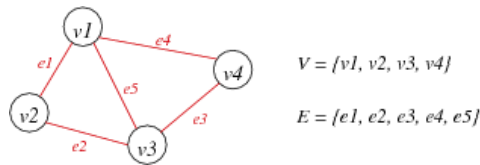
- is a subgraph of G containing all of V
- and is a single tree (connected, **no cycles**)

A *spanning forest* of non-connected graph $G = (V, E)$

- is a subgraph of G containing all of V
- and is a set of trees (not connected, no cycles),
 - with one tree for each *connected component*

Exercise #2: Graph Terminology

15/67



1. How many edges to remove to obtain a spanning tree?
2. How many different spanning trees?

1. 2
2. $\frac{5 \cdot 4}{2} - 2 = 8$ spanning trees (no spanning tree if we remove $\{e1, e2\}$ or $\{e3, e4\}$)

... Graph Terminology

17/67

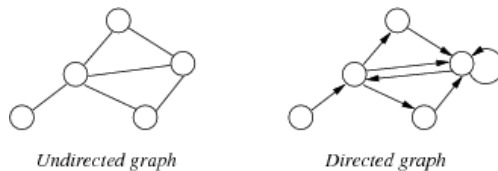
Undirected graph

- $edge(u, v) = edge(v, u)$, no self-loops (i.e. no $edge(v, v)$)

Directed graph

- $edge(u, v) \neq edge(v, u)$, can have self-loops (i.e. $edge(v, v)$)

Examples:



... Graph Terminology

Other types of graphs ...

Weighted graph

- each edge has an associated value (weight)
- e.g. road map (weights on edges are distances between cities)

Multi-graph

- allow multiple edges between two vertices
- e.g. function call graph ($f()$ calls $g()$ in several places)

Graph Data Structures


Graph Representations

20/67

Defining graphs:

- need some way of identifying vertices
- could give diagram showing edges and vertices
- could give a list of edges

E.g. four representations of the same graph:



... Graph Representations

21/67

We will discuss three different graph data structures:

1. Array of edges
2. Adjacency matrix
3. Adjacency list


Array-of-edges Representation

22/67


Edges are represented as an array of Edge values (= pairs of vertices)

- space efficient representation

- adding and deleting edges is slightly complex
- undirected: order of vertices in an Edge doesn't matter
- directed: order of vertices in an Edge encodes direction



[(0,1), (1,2), (1,3), (2,3)]



[(1,0), (1,1), (0,2), (0,3), (2,3)]

For simplicity, we always assume vertices to be numbered $0 \dots V-1$

... Array-of-edges Representation

23/67

Graph initialisation

```
newGraph(V):
|   Input   number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate enough memory for g.edges[]
|   return g
```

... Array-of-edges Representation

24/67

Edge insertion

```
insertEdge(g, (v,w)):
|   Input graph g, edge (v,w)
|
|   i=0
|   while i<g.nE ^ (v,w)≠g.edges[i] do
|       i=i+1
|   end while
|   if i=g.nE then           // (v,w) not found
|       g.edges[i]=(v,w)
|       g.nE=g.nE+1
|   end if
```

... Array-of-edges Representation

25/67

Edge removal

```
removeEdge(g, (v,w)):
|   Input graph g, edge (v,w)
|
|   i=0
```

```
while i<g.nE ^ (v,w)≠g.edges[i] do
    i=i+1
end while
if i<g.nE then           // (v,w) found
    g.edges[i]=g.edges[g.nE-1] // replace by last edge in array
    g.nE=g.nE-1
end if
```

Cost Analysis

26/67

Storage cost: $O(E)$

Cost of operations:

- initialisation: $O(1)$
- insert edge: $O(E)$ (assuming edge array has space)
- delete edge: $O(E)$ (need to find edge in edge array)

If array is full on insert

- allocate space for a bigger array, copy edges across \Rightarrow still $O(E)$

If we maintain edges in order

- use binary search to find edge $\Rightarrow O(\log E)$

Exercise #3: Array-of-edges Representation

27/67

Assuming an array-of-edges representation ...

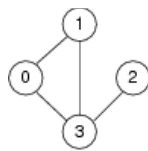
Write an algorithm to output all edges of the graph

```
show(g):
|   Input graph g
|
|   for all i=0 to g.nE-1 do
|       print g.edges[i]
|   end for
```

Adjacency Matrix Representation

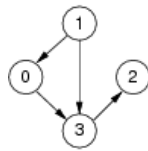
29/67

Edges represented by a $V \times V$ matrix



Undirected graph

A	0	1	2	3
0	0	1	0	1
1	1	0	0	1
2	0	0	0	1
3	1	1	1	0



Directed graph

A	0	1	2	3
0	0	0	0	1
1	1	0	0	1
2	0	0	0	0
3	0	0	1	0

... Adjacency Matrix Representation

30/67

Advantages

- easily implemented as 2-dimensional array
- can represent graphs, digraphs and weighted graphs
 - graphs: symmetric boolean matrix
 - digraphs: non-symmetric boolean matrix
 - weighted: non-symmetric matrix of weight values

Disadvantages:

- if few edges (sparse) \Rightarrow memory-inefficient

... Adjacency Matrix Representation

31/67

Graph initialisation

```
newGraph(V):
|   Input number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate memory for g.edges[][]
|   for all i,j=0..V-1 do
|       g.edges[i][j]=0    // false
|   end for
|   return g
```

... Adjacency Matrix Representation

32/67

Edge insertion

```
insertEdge(g, (v,w)):
```

```
Input graph g, edge (v,w)
```

```
if g.edges[v][w]=0 then // (v,w) not in graph
|   g.edges[v][w]=1      // set to true
|   g.edges[w][v]=1
|   g.nE=g.nE+1
end if
```

... Adjacency Matrix Representation

33/67

Edge removal

```
removeEdge(g, (v,w)):
```

```
Input graph g, edge (v,w)
```

```
if g.edges[v][w]≠0 then // (v,w) in graph
|   g.edges[v][w]=0      // set to false
|   g.edges[w][v]=0
|   g.nE=g.nE-1
end if
```

Exercise #4: Show Graph

34/67

Assuming an adjacency matrix representation ...

Write an algorithm to output all edges of the graph (no duplicates!)

... Adjacency Matrix Representation

35/67

```
show(g):
|   Input graph g
|
|   for all i=0 to g.nV-1 do
|       for all j=i+1 to g.nV-1 do
|           if g.edges[i][j]≠0 then
|               print i-"-"j
|           end if
|       end for
|   end for
```

Exercise #5:

36/67

Analyse storage cost and time complexity of adjacency matrix representation

Storage cost: $O(V^2)$

If the graph is sparse, most storage is wasted.


Cost of operations:

- initialisation: $O(V^2)$ (initialise $V \times V$ matrix)
- insert edge: $O(1)$ (set two cells in matrix)
- delete edge: $O(1)$ (unset two cells in matrix)

... Adjacency Matrix Representation

38/67

A storage optimisation: store only top-right part of matrix.




New storage cost: $V-1$ int ptrs + $V(V+1)/2$ ints (but still $O(V^2)$)

Requires us to always use edges (v,w) such that $v < w$.

Adjacency List Representation

39/67

For each vertex, store linked list of adjacent vertices:



... Adjacency List Representation

40/67

Advantages

- relatively easy to implement in languages like C
- can represent graphs and digraphs
- memory efficient if $E:V$ relatively small

Disadvantages:

- one graph has many possible representations
(unless lists are ordered by same criterion e.g. ascending)

... Adjacency List Representation

41/67

Graph initialisation

```
newGraph(V):
    Input number of nodes V
    Output new empty graph

    g.nV = V    // #vertices (numbered 0..V-1)
    g.nE = 0    // #edges
    allocate memory for g.edges[]
    for all i=0..V-1 do
        g.edges[i]=NULL    // empty list
    end for
    return g
```

... Adjacency List Representation

42/67

Edge insertion:

```
insertEdge(g, (v,w)):
    Input graph g, edge (v,w)

    if !inLL(g.edges[v],w) then    // (v,w) not in graph
        insertLL(g.edges[v],w)
        insertLL(g.edges[w],v)
        g.nE=g.nE+1
    end if
```

... Adjacency List Representation

43/67

Edge removal:

```
removeEdge(g, (v,w)):
    Input graph g, edge (v,w)

    if inLL(g.edges[v],w) then    // (v,w) in graph
        deleteLL(g.edges[v],w)
        deleteLL(g.edges[w],v)
        g.nE=g.nE-1
    end if
```

Exercise #6:

44/67

Storage cost: $O(E)$

Cost of operations:

- initialisation: $O(V)$ (initialise V lists)
- insert edge: $O(1)$ (insert one vertex into list)
- delete edge: $O(E)$ (need to find vertex in list)

If vertex lists are sorted

- insert requires search of list $\Rightarrow O(E)$
- delete always requires a search, regardless of list order

Comparison of Graph Representations

46/67

	array of edges	adjacency matrix	adjacency list
space usage	E	V^2	$V+E$
initialise	I	V^2	V
insert edge	E	I	I
remove edge	E	I	E

Other operations:

	array of edges	adjacency matrix	adjacency list
disconnected(v)?	E	V	I
isPath(x,y)?	$E \cdot \log V$	V^2	$V+E$
copy graph	E	V^2	$V+E$
destroy graph	I	V	$V+E$

Graph Abstract Data Type

Graph ADT

48/67

Data:

- set of edges, set of vertices

Operations:

- building: create graph, add edge
- deleting: remove edge, drop whole graph
- scanning: check if graph contains a given edge

Things to note:

- set of vertices is fixed when graph initialised
- we treat vertices as ints, but could be arbitrary Items

... Graph ADT

49/67

Graph ADT interface **graph.h**

```
// graph representation is hidden
typedef struct GraphRep *Graph;

// vertices denoted by integers 0..N-1
typedef int Vertex;

// edges are pairs of vertices (end-points)
typedef struct Edge { Vertex v; Vertex w; } Edge;

// operations on graphs
Graph newGraph(int V); // new graph with V vertices
void insertEdge(Graph, Edge);
void removeEdge(Graph, Edge);
bool adjacent(Graph, Vertex, Vertex); /* is there an edge
                                         between two vertices */


void freeGraph(Graph);
```

Graph ADT (Array of Edges)

50/67

Implementation of GraphRep (array-of-edges representation)

```
typedef struct GraphRep {
    Edge *edges; // array of edges
    int nV; // #vertices (numbered 0..nV-1)
    int nE; // #edges
    int n; // size of edge array
} GraphRep;
```



... Graph ADT (Array of Edges)

51/67

Implementation of graph initialisation (array-of-edges representation)

```
Graph newGraph(int V) {
    assert(V >= 0);
    Graph g = malloc(sizeof(GraphRep));    assert(g != NULL);

    g->nV = V; g->nE = 0;
    // allocate enough memory for edges
    g->n = Enough;
    g->edges = malloc(g->n*sizeof(Edge));    assert(g->edges != NULL);

    return g;
}
```

How much is enough? ... No more than $V(V-1)/2$... Much less in practice (sparse graph)

... Graph ADT (Array of Edges)

52/67

Implementation of edge insertion/removal (array-of-edges representation)

```
// check if two edges are equal
bool eq(Edge e1, Edge e2) {
    return ( (e1.v == e2.v && e1.w == e2.w)
            || (e1.v == e2.w && e1.w == e2.v) );
}

void insertEdge(Graph g, Edge e) {
    // ensure that g exists and array of edges isn't full
    assert(g != NULL && g->nE < g->n);
    int i = 0;
    while (i < g->nE && !eq(e, g->edges[i]))
        i++;
    if (i == g->nE)                // edge e not found
        g->edges[g->nE++] = e;
}

void removeEdge(Graph g, Edge e) {
    assert(g != NULL);            // ensure that g exists
    int i = 0;
    while (i < g->nE && !eq(e, g->edges[i]))
        i++;
    if (i < g->nE)                // edge e found
        g->edges[i] = g->edges[--g->nE];
}
```

Exercise #7: Checking Neighbours (i)

53/67

Assuming an array-of-edges representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```


```
bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL);
    Edge e;
    e.v = x; e.w = y;
    int i = 0;
    while (i < g->nE) {
        if (eq(e, g->edges[i]))    // edge found
            return true;
        i++;
    }
    return false;                // edge not found
}
```

Graph ADT (Adjacency Matrix)

55/67

Implementation of GraphRep (adjacency-matrix representation)

```
typedef struct GraphRep {
    int **edges; // adjacency matrix
    int nV;      // #vertices
    int nE;      // #edges
} GraphRep;
```



... Graph ADT (Adjacency Matrix)

56/67

Implementation of graph initialisation (adjacency-matrix representation)

```
Graph newGraph(int V) {
    assert(V >= 0);
    int i;

    Graph g = malloc(sizeof(GraphRep));    assert(g != NULL);
    g->nV = V; g->nE = 0;

    // allocate memory for each row
    g->edges = malloc(V * sizeof(int *));    assert(g->edges != NULL);
    // allocate memory for each column and initialise with 0
    for (i = 0; i < V; i++) {
        g->edges[i] = calloc(V, sizeof(int));    assert(g->edges[i] != NULL);
    }
    return g;
}
```


standard library function `calloc(size_t nelems, size_t nbytes)`

- allocates a memory block of size `nelems*nbytes`
- and sets all bytes in that block to *zero*

... Graph ADT (Adjacency Matrix)

57/67

Implementation of edge insertion/removal (adjacency-matrix representation)

```
// check if vertex is valid in a graph
bool validV(Graph g, Vertex v) {
    return (g != NULL && v >= 0 && v < g->nV);
}

void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (!g->edges[e.v][e.w]) { // edge e not in graph
        g->edges[e.v][e.w] = 1;
        g->edges[e.w][e.v] = 1;
        g->nE++;
    }
}

void removeEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (g->edges[e.v][e.w]) { // edge e in graph
        g->edges[e.v][e.w] = 0;
        g->edges[e.w][e.v] = 0;
        g->nE--;
    }
}
```

Exercise #8: Checking Neighbours (ii)

58/67

Assuming an adjacency-matrix representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```
bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL && validV(g,x) && validV(g,y));

    return (g->edges[x][y] != 0);
}
```


Graph ADT (Adjacency List)

60/67

Implementation of GraphRep (adjacency-list representation)

```
typedef struct GraphRep {
    Node **edges; // array of lists
    int    nV;    // #vertices
    int    nE;    // #edges
} GraphRep;
```

```
typedef struct Node {
    Vertex    v;
    struct Node *next;
} Node;
```



... Graph ADT (Adjacency List)

61/67

Implementation of graph initialisation (adjacency-list representation)

```
Graph newGraph(int V) {
    assert(V >= 0);
    int i;

    Graph g = malloc(sizeof(GraphRep));
    g->nV = V; g->nE = 0;

    // allocate memory for array of lists
    g->edges = malloc(V * sizeof(Node *));
    assert(g->edges != NULL);
    for (i = 0; i < V; i++)
        g->edges[i] = NULL;

    return g;
}
```

... Graph ADT (Adjacency List)

62/67

Implementation of edge insertion/removal (adjacency-list representation)

```
void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (!inLL(g->edges[e.v], e.w)) { // edge e not in graph
        g->edges[e.v] = insertLL(g->edges[e.v], e.w);
        g->edges[e.w] = insertLL(g->edges[e.w], e.v);
        g->nE++;
    }
}
```

```
void removeEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (inLL(g->edges[e.v], e.w)) { // edge e in graph
        g->edges[e.v] = deleteLL(g->edges[e.v], e.w);
        g->edges[e.w] = deleteLL(g->edges[e.w], e.v);
        g->nE--;
    }
}
```

inLL, insertLL, deleteLL are standard linked list operations (as discussed in week 3)

Exercise #9: Checking Neighbours (iii)

63/67

Assuming an adjacency list representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```
bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL && validV(g,x));


    return inLL(g->edges[x], y);
}
```

Exercise #10: Graph ADT Client

65/67

Write a program that uses the graph ADT to

- build the graph depicted below
- print all the nodes that are incident to vertex 1 in ascending order



```
#include <stdio.h>
#include "Graph.h"
```

```
#define NODES 4
#define NODE_OF_INTEREST 1
```

```
int main(void) {
    Graph g = newGraph(NODES);
```

```
Edge e;
e.v = 0; e.w = 1; insertEdge(g,e);
e.v = 0; e.w = 3; insertEdge(g,e);
e.v = 1; e.w = 3; insertEdge(g,e);
e.v = 3; e.w = 2; insertEdge(g,e);

int v;
for (v = 0; v < NODES; v++) {
    if (adjacent(g, v, NODE_OF_INTEREST))
        printf("%d\n", v);
}

freeGraph(g);
return 0;
}
```

Summary

67/67

- Graph terminology
 - vertices, edges, vertex degree, connected graph, tree
 - path, cycle, clique, spanning tree, spanning forest
- Graph representations
 - array of edges
 - adjacency matrix
 - adjacency lists
- Suggested reading:
 - Sedgewick, Ch.17.1-17.5

Produced: 29 Aug 2017