

## Week 11: Search Tree Algorithms 2

### Tree Review

1/67

Binary search trees ...

- data structures designed for  $O(\log n)$  search
- consist of nodes containing item (incl. key) and two links
- can be viewed as recursive data structure (subtrees)
- have overall ordering (data(Left) < root < data(Right))
- insert new nodes as leaves (or as root), delete from anywhere
- have structure determined by insertion order (*worst*:  $O(n)$ )
- operations: insert, delete, search, rotate, rebalance, ...

### Randomised BST Insertion

2/67

Effects of order of insertion on BST shape:

- best case (for at-leaf insertion): keys inserted in pre-order (median key first, then median of lower half, median of upper half, etc.)
- worst case: keys inserted in ascending/descending order
- average case: keys inserted in *random* order  $\Rightarrow O(\log_2 n)$

Tree ADT has no control over order that keys are supplied.

Can the algorithm itself introduce some *randomness*?

In the hope that this randomness helps to balance the tree ...

### Sidetrack: Random Numbers

3/67

How can a computer pick a number at random?

- it cannot

Software can only produce *pseudo random numbers*.

- a pseudo random number is one that is predictable
  - (although it may appear unpredictable)

$\Rightarrow$  Implementation may deviate from expected theoretical behaviour

### ... Sidetrack: Random Numbers

4/67

The most widely-used technique is called the *Linear Congruential Generator (LCG)*

- it uses a **recurrence** relation:

- $X_{n+1} = (a \cdot X_n + c) \bmod m$ , where:
  - $m$  is the "modulus"
  - $a, 0 < a < m$  is the "multiplier"
  - $c, 0 \leq c \leq m$  is the "increment"
  - $X_0$  is the "*seed*"
- if  $c=0$  it is called a *multiplicative congruential generator*

LCG is not good for applications that need extremely high-quality random numbers

- the period length is too short (length of the sequence at which point it repeats itself)
- a short period means the numbers are correlated

### ... Sidetrack: Random Numbers

5/67

Trivial example:

- for simplicity assume  $c=0$
- so the formula is  $X_{n+1} = a \cdot X_n \bmod m$
- try  $a=11, m=31$ , which generates the sequence:

11, 28, 29, 9, 6, 4, 13, 19, 23, 5, 24, 16, 21, 14, 30, 20, 3, 2, 22, 25, 27, 18, 12, 8, 26, 7, 15, 10, 17, 1, 11, 28, 29, 9, 6, 4, 13, 19, 23, 5, 24, 16, 21, 14, 30, 20, 3, 2, 22, 25, 27, 18, 12, 8, 26, 7, 15, 10, 17, 1, 11, 28, 29, 9, 6, 4, 13, 19, 23, 5, 24, 16, 21, 14, 30, 20, 3, 2, 22, 25, 27, 18, 12, 8, 26, 7, 15, 10, 17, 1, 11, 28, 29, 9, 6, 4, 13, 19, 23, 5, 24, 16, 21, 14, 30, 20, 3, 2, 22, 25, 27, 18, 12, 8, 26, 7, 15, 10, 17, 1, 11, 28, 29, 9, 6, 4, 13, 19, 23, 5, 24, 16, 21, 14, 30, 20, 3, 2, 22, 25, 27, 18, 12, 8, 26, 7, 15, 10, 17, 1, ...

- all the integers from 1 to 30 are here

### ... Sidetrack: Random Numbers

6/67

Another trivial example:

- again let  $c=0$
- try  $a=12, m=30$ 
  - that is,  $X_{n+1} = 12 \cdot X_n \bmod 30$
  - which generates the sequence:

12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6, ...

- notice the period length ... clearly a terrible sequence

### ... Sidetrack: Random Numbers

7/67

It is a complex task to pick good numbers. A bit of history:

Lewis, Goodman and Miller (1969) suggested

- $X_{n+1} = 7^5 \cdot X_n \bmod (2^{31}-1)$
- note:
  - $7^5$  is 16807
  - $2^{31}-1$  is 2147483647
  - $X_0 = 0$  is not a good seed value

Most compilers use LCG-based algorithms that are slightly more involved; see [www.mscs.dal.ca/~selinger/random/](http://www.mscs.dal.ca/~selinger/random/) for details (including a short C program that produces the exact same pseudo-random numbers as `gcc` for any given seed value)

## ... Sidetrack: Random Numbers

8/67

- Two functions are required:

```
srandom(int seed) // sets its argument as the seed
```

```
random() // uses a LCG technique to generate random
          // numbers in the range 0 .. RAND_MAX
```

where the constant `RAND_MAX` is defined in `stdlib.h`  
(depends on the computer: on the CSE network, `RAND_MAX` = 2147483647)

- The period length of this random number generator is very large  
approximately  $16 \cdot ((2^{31}) - 1)$

## ... Sidetrack: Random Numbers

9/67

To convert the return value of `random()` to a number between 0 .. RANGE

- compute the remainder after division by `RANGE+1`

Using the remainder to compute a random number is not the best way:

- can generate a 'better' random number by using a more complex division
- but good enough for most purposes

Some applications require more sophisticated, *cryptographically secure* pseudo random numbers

## Exercise #1: Random Numbers

10/67

Write a program to simulate 10,000 rounds of Two-up.

- Assume a \$10 bet at each round
- Compute the overall outcome and average per round

```
#include <stdlib.h>
#include <stdio.h>
```

```
#define RUNS 10000
#define BET 10
```

```
int main(void) {
    int coin1, coin2, n, sum = 0;
    for (n = 0; n < RUNS; n++) {
        do {
            coin1 = random() % 2;
            coin2 = random() % 2;
        } while (coin1 != coin2);
        if (coin1==1 && coin2==1)
            sum += BET;
        else
            sum -= BET;
    }
    printf("Final result: %d\n", sum);
    printf("Average outcome: %f\n", (float) sum / RUNS);
    return 0;
}
```

## ... Sidetrack: Random Numbers

12/67

### Seeding

There is one significant problem:

- every time you run a program with the same seed, you get exactly the same sequence of 'random' numbers (why?)

To vary the output, can give the random seeder a starting point that varies with time

- an example of such a starting point is the current time, *time(NULL)*  
(NB: this is different from the UNIX command `time`, used to measure program running time)

```
#include <time.h>
time(NULL) // returns the time as the number of seconds
            // since the Epoch, 1970-01-01 00:00:00 +0000
```

```
// time(NULL) on October 10th, 2017, 12:59pm was 1507600763
// time(NULL) about a minute later was 1507600825
```

## Randomised BST Insertion

13/67

Approach: normally do leaf insert, randomly do root insert.

```
insertRandom(tree,item)
|   Input   tree, item
|   Output tree with item randomly inserted
|
|   if tree is empty then
|       return new node containing item
|   end if
|   // p/q chance of doing root insert
|   if random() mod q < p then
```

```

return insertAtRoot(tree,item)
else
return insertAtLeaf(tree,item)
end if

```

E.g. 30% chance  $\Rightarrow$  choose  $p=3, q=10$

## ... Randomised BST Insertion

14/67

Cost analysis:

- similar to cost for inserting keys in random order:  $O(\log_2 n)$
- does not rely on keys being supplied in random order

Approach can also be applied to deletion:

- standard method promotes inorder successor to root
- for the randomised method ...
  - promote inorder successor from right subtree, OR
  - promote inorder predecessor from left subtree

## Splay Trees

16/67

A kind of "self-balancing" tree ...

Splay tree insertion modifies insertion-at-root method:

- by considering *parent-child-grandchild* (three level analysis)
- by performing double-rotations based on p-c-g orientation

The idea: appropriate double-rotations improve tree balance.

Splay tree implementations also do *rotation-in-search*:


- can provide similar effect to periodic rebalance
- improves balance, but makes search more expensive

## ... Splay Trees

17/67

Cases for splay tree double-rotations:


- case 1: grandchild is left-child of left-child
- case 2: grandchild is right-child of left-child
- case 3: grandchild is left-child of right-child
- case 4: grandchild is right-child of right-child



## ... Splay Trees

18/67


Example: double-rotation case for left-child of left-child:



## ... Splay Trees

19/67

Example: double-rotation case for right-child of left-child:



## ... Splay Trees

20/67

Algorithm for splay tree insertion:


```
insertSplay(tree,item):
```

**Input** tree, item  
**Output** tree with item splay-inserted

```

if tree is empty then return new node containing item
else if item=data(tree) then return tree
else if item<data(tree) then
  if left(tree) is empty then
    left(tree)=new node containing item
  else if item<data(left(tree)) then
    // Case 1: left-child of left-child
    left(left(tree))=insertSplay(left(left(tree)),item)
    left(tree)=rotateRight(left(tree))
    tree=rotateRight(tree)
  else // Case 2: right-child of left-child
    right(left(tree))=insertSplay(right(left(tree)),item)
    left(tree)=rotateLeft(left(tree))
  end if
  return rotateRight(tree)
else if item>data(tree) then
  if right(tree) is empty then
    right(tree)=new node containing item
  else if item<data(right(tree)) then
    // Case 3: left-child of right-child
    left(right(tree))=insertSplay(left(right(tree)),item)
    right(tree)=rotateRight(right(tree))
  else // Case 4: right-child of right-child
    right(right(tree))=insertSplay(right(right(tree)),item)
    right(tree)=rotateLeft(right(tree))
    tree=rotateLeft(tree)
  end if
  return rotateLeft(tree)
end if

```



## ... Splay Trees

23/67

Searching in splay trees:

```

searchSplay(tree,item):
  Input tree, item
  Output address of item if found in tree
  NULL otherwise

  if tree=NULL then
    return NULL
  else
    tree=splay(tree,item)
    if data(tree)=item then
      return tree
    else
      return NULL
    end if
  end if


```

where **splay()** is similar to **insertSplay()**,  
except that it doesn't add a node ... simply moves **item** to root if found, or nearest node if not found

## Exercise #2: Splay Trees

21/67

Insert 18 into this splay tree:




Correction:


## Exercise #3: Splay Trees

24/67

If we search for 22 in the splay tree



... how does this affect the tree?



## ... Splay Trees

26/67

Analysis of splay tree performance:

- assume that we "splay" for both insert and search
- consider:  $m$  insert+search operations,  $n$  nodes
- total number of comparisons: average  $O((n+m) \cdot \log_2(n+m))$

Gives good overall (amortized) cost.

- insert cost not significantly different to insert-at-root
- search cost increases, but ...
  - improves balance on each search
  - moves frequently accessed nodes closer to root

But ... still has worst-case search cost  $O(n)$

## Real Balanced Trees

## Better Balanced Binary Search Trees

28/67

So far, we have seen ...

- randomised trees ... make poor performance unlikely
- occasional rebalance ... fix balance periodically
- splay trees ... reasonable amortized performance
- but both types still have  $O(n)$  worst case

Ideally, we want both average/worst case to be  $O(\log n)$

- AVL trees ... fix imbalances as soon as they occur
- 2-3-4 trees ... use varying-sized nodes to assist balance
- red-black trees ... isomorphic to 2-3-4, but binary nodes

## AVL Trees

## AVL Trees

Invented by Georgy Adelson-Velsky and Evgenii Landis

Approach:

- insertion (at leaves) may cause imbalance
- repair balance as soon as we notice imbalance
- repairs done locally, not by overall tree restructure

A tree is unbalanced when:  $\text{abs}(\text{height}(\text{left}) - \text{height}(\text{right})) > 1$

This can be repaired by a single rotation:

- if left subtree too deep, rotate right
- if right subtree too deep, rotate left

Problem: determining height/depth of subtrees may be expensive.

## ... AVL Trees

31/67

Implementation of AVL insertion

```

insertAVL(tree,item):
  Input  tree, item
  Output tree with item AVL-inserted



  if tree is empty then
    return new node containing item
  else if item=data(tree) then
    return tree
  else
    if item<data(tree) then
      left(tree)=insertAVL(left(tree),item)
    else if item>data(tree) then
      right(tree)=insertAVL(right(tree),item)
    end if
    if height(left(tree))-height(right(tree)) > 1 then
      if item>data(left(tree)) then
        left(tree)=rotateLeft(left(tree))
      end if
      tree=rotateRight(tree)
    else if height(right(tree))-height(left(tree)) > 1 then
      if item<data(right(tree)) then
        right(tree)=rotateRight(right(tree))
      end if
      tree=rotateLeft(tree)
    end if
    return tree
  end if
  
```

30/67

## Exercise #4: AVL Trees

32/67

Insert 27 into the AVL tree




## ... AVL Trees

34/67

Analysis of AVL trees:

- trees are *height*-balanced; subtree depths differ by  $\pm 1$
- average/worst-case search performance of  $O(\log n)$
- *require* extra data to be stored in each node (efficiency)
- may not be *weight*-balanced; subtree sizes may differ




## 2-3-4 Trees

## 2-3-4 Trees

36/67

2-3-4 trees have three kinds of nodes

- 2-nodes, with two children (same as normal BSTs)
- 3-nodes, two values and three children
- 4-nodes, three values and four children



## ... 2-3-4 Trees

37/67

2-3-4 trees are ordered similarly to BSTs



In a *balanced* 2-3-4 tree:

- all leaves are at same distance from the root

2-3-4 trees grow "upwards" from the leaves.

## ... 2-3-4 Trees

38/67

Possible 2-3-4 tree data structure:

```
typedef struct node {
    int      order;      // 2, 3 or 4
    int      data[3];    // items in node
    struct node *child[4]; // links to subtrees
} node;
```

## ... 2-3-4 Trees

39/67

Searching in 2-3-4 trees:

```
Search(tree,item):
    Input  tree, item
    Output address of item if found in 2-3-4 tree
           NULL otherwise

    if tree is empty then
        return NULL
    else
        i=0
        while i<tree.order-1 ^ item>tree.data[i] do
            i=i+1 // find relevant slot in data[]
        end while
        if item=tree.data[i] then // item found
            return address of tree.data[i]
```

```

else      // keep looking in relevant subtree
  return Search(tree.child[i],item)
end if
end if

```

## ... 2-3-4 Trees

40/67

2-3-4 tree searching cost analysis:

- as for other trees, worst case determined by height  $h$
- 2-3-4 trees are always balanced  $\Rightarrow$  height is  $O(\log n)$
- worst case for height: all nodes are 2-nodes  
same case as for balanced BSTs, i.e.  $h \approx \log_2 n$
- best case for height: all nodes are 4-nodes  
balanced tree with branching factor 4, i.e.  $h \approx \log_4 n$

## Insertion into 2-3-4 Trees

41/67


Insertion algorithm:

- find leaf node where Item belongs (via search)
- if not full (i.e. order  $< 4$ )
  - insert Item in this node, order++
- if node is full (i.e. contains 3 items)
  - split into two 2-nodes as leaves
  - promote middle element to parent
  - insert item into appropriate leaf 2-node
  - if parent is a 4-node
    - continue split/promote upwards
  - if promote to root, and root is a 4-node
    - split root node and add new root

## ... Insertion into 2-3-4 Trees

42/67

Insertion into a 2-node or 3-node:




Insertion into a 4-node (requires a split):



## ... Insertion into 2-3-4 Trees

43/67


Splitting the root:



## ... Insertion into 2-3-4 Trees

44/67



Building a 2-3-4 tree ... 7 insertions:



## Exercise #5: Insertion into 2-3-4 Tree

45/67

Show what happens when D, S, F, U are inserted into this tree:



Insertion algorithm:

```
insert(tree,item):
    Input  2-3-4 tree, item
    Output tree with item inserted

    if tree is empty then
        return new node containing item
    end if
    node=Search(tree,item)
    parent=parent of node
    if node.order<4 then
        insert item into node
        increment node.order
    else
        promote = node.data[1]      // middle value
        nodeL   = new node containing data[0]
        nodeR   = new node containing data[2]
        if item<node.data[1] then
            insert(nodeL,item)
        else
            insert(nodeR,item)
        end if
        insert(parent,promote)
        while parent.order=4 do
            continue promote/split upwards
        end while
        if parent is root ^ parent.order=4 then
            split root, making new root
        end if
    end if
end if
```

Variations on 2-3-4 trees ...

Variation #1: why stop at 4? why not 2-3-4-5 trees? or  $M$ -way trees?

- allow nodes to hold up to  $M-1$  items, and at least  $M/2$
- if each node is a disk-page, then we have a *B-tree* (databases)
- for B-trees, depending on item size,  $M > 100/200/400$

Variation #2: don't have "variable-sized" nodes

- use standard BST nodes, augmented with one extra piece of data
- implement similar strategy as 2-3-4 trees  $\rightarrow$  red-black trees.

## Red-Black Trees

## Red-Black Trees

*Red-black trees* are a representation of 2-3-4 trees using BST nodes.

- each node needs one extra value to encode link type
- but we no longer have to deal with different kinds of nodes

Link types:

- *red* links ... combine nodes to represent 3- and 4-nodes
- *black* links ... analogous to "ordinary" BST links (child links)

Advantages:

- standard BST search procedure works unmodified
- get benefits of 2-3-4 tree self-balancing (although deeper)

## Red-Black Trees

Definition of a *red-black tree*

- a BST in which each node is marked red or black
- no two red nodes appear consecutively on any path
- a red node corresponds to a 2-3-4 sibling of its parent
- a black node corresponds to a 2-3-4 child of its parent


*Balanced* red-black tree

- all paths from root to leaf have same number of black nodes

Insertion algorithm: avoids worst case  $O(n)$  behaviour


Search algorithm: standard BST search

Representing 4-nodes in red-black trees:



Some texts colour the links rather than the nodes.






## ... Red-Black Trees

54/67

Equivalent trees (one 2-3-4, one red-black):



## ... Red-Black Trees

55/67

Red-black tree implementation:

```
typedef enum {RED, BLACK} Colour;
typedef struct node *RBTree;
typedef struct node {
    int data; // actual data
    Colour colour; // relationship to parent
    RBTree left; // left subtree
    RBTree right; // right subtree
} node;

#define colour(tree) ((tree)->colour)
#define isRed(tree) ((tree) != NULL && (tree)->colour == RED)
```

RED = node is part of the same 2-3-4 node as its parent (sibling)

BLACK = node is a child of the 2-3-4 node containing the parent

## ... Red-Black Trees

56/67

New nodes are always red:


```
RBTree newNode(Item it) {
    RBTree new = malloc(sizeof(Node));
    assert(new != NULL);
    data(new) = it;
    colour(new) = RED;
    left(new) = right(new) = NULL;
    return new;
}
```

## ... Red-Black Trees

57/67

Node.red allows us to distinguish links

- black = parent node is a "real" parent
- red = parent node is a 2-3-4 neighbour



## ... Red-Black Trees

58/67

Search method is standard BST search:

```
SearchRedBlack(tree, item):
    Input tree, item
    Output true if item found in red-black tree
           false otherwise

    if tree is empty then
        return false
    else if item < data(tree) then
        return SearchRedBlack(left(tree), item)
    else if item > data(tree) then
        return SearchRedBlack(right(tree), item)
    else // found
        return true
    end if
```

## Red-Black Tree Insertion

59/67

Insertion is more complex than for standard BSTs

- need to recall direction of last branch (L or R)
- need to recall whether parent link is red or black
- splitting/promoting implemented by rotateLeft/rotateRight
- several cases to consider depending on colour/direction combinations

### ... Red-Black Tree Insertion

60/67

High-level description of insertion algorithm:

```
insertRB(tree,item,inRight):
  Input tree, item, inRight indicating direction of last branch
  Output tree with it inserted

  if tree is empty then
    return newNode(item)
  end if
  if left(tree) and right(tree) both are RED then
    split 4-node
  end if
  recursive insert cases (cf. regular BST)
  re-arrange links/colours after insert
  return modified tree
```


```
insertRedBlack(tree,item):
  Input red-black tree, item
  Output tree with item inserted

  tree=insertRB(tree,item,false)
  colour(tree)=BLACK
  return tree
```

### ... Red-Black Tree Insertion

61/67

Splitting a 4-node, in a red-black tree:



Algorithm:

```
if isRed(left(currentTree)) ^ isRed(right(currentTree)) then
  colour(currentTree)=RED
  colour(left(currentTree))=BLACK
  colour(right(currentTree))=BLACK
end if
```

### ... Red-Black Tree Insertion

62/67

Simple recursive insert (a la BST):



Algorithm:


```
if item < data(tree) then
  left(tree)=insertRB(left(tree),item,false)
  re-arrange links/colours after insert
else // item larger than data in root
  right(tree)=insertRB(right(tree),item,true)
  re-arrange links/colours after insert
end if
```

Not affected by colour of tree node.

### ... Red-Black Tree Insertion

63/67

Re-arrange after insert (1): two successive red links = newly-created 4-node



Algorithm:

```
if isRed(left(currentTree)) ^ isRed(left(left(currentTree))) then
  currentTree=rotateRight(currentTree)
  colour(currentTree)=BLACK
  colour(right(currentTree))=RED
end if
```

### ... Red-Black Tree Insertion

64/67

Re-arrange after insert (2): "normalise" direction of successive red links



Algorithm:

```
if inRight  $\wedge$  isRed(currentTree)  $\wedge$  isRed(left(currentTree)) then  
    currentTree=rotateRight(currentTree)  
end if
```


---

## ... Red-Black Tree Insertion

65/67

Example of insertion, starting from empty tree:

22, 12, 8, 15, 11, 19, 43, 44, 45, 42, 41, 40, 39



## Red-black Tree Performance

66/67

Cost analysis for red-black trees:

- tree is well-balanced; worst case search is  $O(\log_2 n)$
- insertion affects nodes down one path; max #rotations is  $2 \cdot h$   
(where  $h$  is the height of the tree)

Only disadvantage is complexity of insertion/deletion code.

Note: red-black trees were popularised by Sedgewick.

---

## Summary

67/67

- Random numbers
  - Self-adjusting trees
    - Splay trees
    - AVL trees
    - 2-3-4 trees
    - Red-black trees
  - Suggested reading:
    - Sedgewick, Ch.13.1-13.4
-