

Strings

Strings

A *string* is a sequence of characters.

An *alphabet* Σ is the set of possible characters in strings.

Examples of strings:

- C program
- HTML document
- DNA sequence
- Digitized image

Examples of alphabets:

- ASCII
- Unicode
- {0,1}
- {A,C,G,T}

... Strings

Notation:

- $length(P)$... #characters in P
- λ ... *empty* string ($length(\lambda) = 0$)
- Σ^m ... set of all strings of length m over alphabet Σ
- Σ^* ... set of all strings over alphabet Σ

$v\omega$ denotes the *concatenation* of strings v and ω

Note: $length(v\omega) = length(v) + length(\omega)$ $\lambda\omega = \omega = \omega\lambda$

... Strings

Notation:

- *substring* of P ... any string Q such that $P = vQ\omega$, for some $v, \omega \in \Sigma^*$
- *prefix* of P ... any string Q such that $P = Q\omega$, for some $\omega \in \Sigma^*$
- *suffix* of P ... any string Q such that $P = \omega Q$, for some $\omega \in \Sigma^*$

Exercise #1: Strings

The string **a/a** of length 3 over the ASCII alphabet has

- how many prefixes?
- how many suffixes?
- how many substrings?

- 4 prefixes: " " "a" "a/" "a/a"
- 4 suffixes: "a/a" "/a" "a" ""
- 6 substrings: " " "a" "/" "a/" "/a" "a/a"

Note:

" " means the same as λ (= empty string)

... Strings

ASCII (American Standard Code for Information Interchange)

- Specifies mapping of 128 characters to integers 0..127
- The characters encoded include:
 - upper and lower case English letters: A-Z and a-z
 - digits: 0-9
 - common punctuation symbols
 - special non-printing characters: e.g. *newline* and *space*

Ascii	Char	Ascii	Char	Ascii	Char	Ascii	Char
0	Null	32	Space	64	@	96	`
1	Start of heading	33	!	65	A	97	a
2	Start of text	34	"	66	B	98	b
3	End of text	35	#	67	C	99	c
4	End of transmit	36	\$	68	D	100	d
5	Enquiry	37	%	69	E	101	e
6	Acknowledge	38	&	70	F	102	f
7	Audible bell	39	'	71	G	103	g
8	Backspace	40	(72	H	104	h
9	Horizontal tab	41)	73	I	105	i
10	Line feed	42	*	74	J	106	j
11	Vertical tab	43	+	75	K	107	k
12	Form feed	44	,	76	L	108	l
13	Carriage return	45	-	77	M	109	m
14	Shift in	46	.	78	N	110	n
15	Shift out	47	/	79	O	111	o
16	Data link escape	48	0	80	P	112	p
17	Device control 1	49	1	81	Q	113	q
18	Device control 2	50	2	82	R	114	r
19	Device control 3	51	3	83	S	115	s
20	Device control 4	52	4	84	T	116	t
21	Neg. acknowledge	53	5	85	U	117	u
22	Synchronous idle	54	6	86	V	118	v
23	End trans. block	55	7	87	W	119	w
24	Cancel	56	8	88	X	120	x
25	End of medium	57	9	89	Y	121	y
26	Substitution	58	:	90	Z	122	z
27	Escape	59	;	91	[123	{
28	File separator	60	<	92	\	124	
29	Group separator	61	=	93]	125	}
30	Record separator	62	>	94	^	126	~
31	Unit separator	63	?	95	_	127	Forward del.

... Strings

Reminder:

In C a string is an array of chars containing ASCII codes

- these arrays have an extra element containing a 0
- the extra 0 can also be written '`\0`' (*null character* or *null-terminator*)
- convenient because don't have to track the length of the string

Because strings are so common, C provides convenient syntax:

`char str[] = "hello";` // same as `char str[] = {'h', 'e', 'l', 'l', 'o', '\0'};`

Note: `str[]` will have 6 elements

... Strings

C provides a number of string manipulation functions via `#include <string.h>`, e.g.

```
strlen()    // length of string
strncpy()   // copy one string to another
strncat()   // concatenate two strings
strstr()    // find substring inside string
```

Example:

```
char *strncat(char *dest, char *src, int n)
```


- appends string `src` to the end of `dest` overwriting the `'\0'` at the end of `dest` and adds terminating `'\0'`
- returns start of string `dest`
- will never add more than `n` characters
(If `src` is less than `n` characters long, the remainder of `dest` is filled with `'\0'` characters. Otherwise, `dest` is not null-terminated.)

Pattern Matching

Pattern Matching

11/65

Example (pattern checked *backwards*):



- *Text* ... abacaab
- *Pattern* ... abacab

... Pattern Matching

12/65

Given two strings T (*text*) and P (*pattern*), the *pattern matching problem* consists of finding a substring of T equal to P

Applications:

- Text editors
- Search engines
- Biological research

... Pattern Matching

13/65

Brute-force pattern matching algorithm

- checks for each possible shift of P relative to T
 - until a match is found, or
 - all placements of the pattern have been tried

BruteForceMatch(T, P):

```
Input  text T of length n, pattern P of length m
Output starting index of a substring of T equal to P
        -1 if no such substring exists

for all i=0..n-m do
```

```
    j=0                                // check from left to right
    while j<m ^ T[i+j]=P[j] do         // test ith shift of pattern
        j=j+1
        if j=m then
            return i                    // entire pattern checked
        end if
    end while
end for
return -1                             // no match found
```

Analysis of Brute-force Pattern Matching

14/65

Brute-force pattern matching runs in $O(n \cdot m)$

Examples of worst case (forward checking):

- $T = \text{aaa...ah}$
- $P = \text{aaah}$
- may occur in DNA sequences
- unlikely in English text

Boyer-Moore Algorithm

15/65


The *Boyer-Moore* pattern matching algorithm is based on two heuristics:

- *Looking-glass heuristic*: Compare P with subsequence of T moving *backwards*
- *Character-jump heuristic*: When a mismatch occurs at $T[i]=c$
 - if P contains $c \Rightarrow$ shift P so as to align the **last** occurrence of c in P with $T[i]$
 - otherwise \Rightarrow shift P so as to align $P[0]$ with $T[i+1]$ (a.k.a. "big jump")

... Boyer-Moore Algorithm

16/65

Example:




... Boyer-Moore Algorithm

17/65

Boyer-Moore algorithm preprocesses pattern P and alphabet Σ to build

- *last-occurrence function* L
 - L maps Σ to integers such that $L(c)$ is defined as
 - the largest index i such that $P[i]=c$, or
 - -1 if no such index exists

Example: $\Sigma = \{a, b, c, d\}$, $P = \text{acab}$



$L(c)$	2	3	1	-1
--------	---	---	---	----

- L can be represented by an array indexed by the numeric codes of the characters
- L can be computed in $O(m+s)$ time ($m \dots$ length of pattern, $s \dots$ size of Σ)

... Boyer-Moore Algorithm

18/65

BoyerMooreMatch(T, P, Σ):

Input text T of length n , pattern P of length m , alphabet Σ
Output starting index of a substring of T equal to P
 -1 if no such substring exists

```


L=lastOccurenceFunction(P,Σ)
i=m-1, j=m-1          // start at end of pattern
repeat
  if T[i]=P[j] then
    if j=0 then
      return i          // match found at i
    else
      i=i-1, j=j-1
    end if
  else
    // character-jump
    i=i+m-min(j,1+L[T[i]])
    j=m-1
  end if
until i≥n
return -1              // no match
  
```

- Biggest jump (m characters ahead) occurs when $L[T[i]] = -1$


... Boyer-Moore Algorithm

19/65

Case 1: $j \leq l + L[c]$



Case 2: $l + L[c] < j$




Exercise #2: Boyer-Moore algorithm

20/65

For the alphabet $\Sigma = \{a, b, c, d\}$

1. compute last-occurrence function L for pattern $P = \mathbf{abacab}$
2. trace Boyer-More on P and text $T = \mathbf{abacaabadcabacabaabb}$
 - how many comparisons are needed?

c	a	b	c	d
$L(c)$	4	5	3	-1



13 comparisons in total

... Boyer-Moore Algorithm

22/65

Analysis of Boyer-Moore algorithm:

- Runs in $O(nm+s)$ time
 - $m \dots$ length of pattern $n \dots$ length of text $s \dots$ size of alphabet
- Example of worst case:
 - $T = \mathbf{aaa \dots a}$
 - $P = \mathbf{baaa}$
- Worst case may occur in images and DNA sequences but unlikely in English texts
 \Rightarrow Boyer-Moore significantly faster than brute-force on English text

Knuth-Morris-Pratt Algorithm

23/65

The *Knuth-Morris-Pratt* algorithm ...

- compares the pattern to the text *left-to-right*
- but shifts the pattern more intelligently than the brute-force algorithm

Reminder:

- Q is a *prefix* of $P \dots P = Q\omega$, for some $\omega \in \Sigma^*$
- Q is a *suffix* of $P \dots P = \omega Q$, for some $\omega \in \Sigma^*$


... Knuth-Morris-Pratt Algorithm

24/65

When a mismatch occurs ...

- what is the most we can shift the pattern to avoid redundant comparisons?

- Answer: the largest *prefix* of $P[0..j]$ that is a *suffix* of $P[1..j]$



... Knuth-Morris-Pratt Algorithm

25/65

KMP preprocesses the pattern to find matches of its prefixes with itself

- *Failure function* $F(j)$ defined as
 - the size of the *largest prefix* of $P[0..j]$ that is also a *suffix* of $P[1..j]$
- if mismatch occurs at $P_j \Rightarrow$ advance j to $F(j-1)$

Example: $P = \text{abaaba}$

j	0	1	2	3	4	5
P_j	a	b	a	a	b	a
$F(j)$	0	0	1	1	2	3

$F(j-1)$

... Knuth-Morris-Pratt Algorithm

26/65

KMPMatch(T, P):

Input text T of length n , pattern P of length m
Output starting index of a substring of T equal to P
 -1 if no such substring exists

$F = \text{failureFunction}(P)$

```

i=0, j=0 // start from left
while i<n do
  if T[i]=P[j] then
    if j=m-1 then
      return i-j // match found at i-j
    else
      i=i+1, j=j+1
    end if
  else
    // mismatch at P[j]

```

```


    if j>0 then
      j=F[j-1] // resume comparing P at F[j-1]
    else
      i=i+1
    end if
  end if
end while
return -1 // no match

```

KMP-Algorithm

1. compute failure function F for pattern $P = \text{abacab}$
2. trace Knuth-Morris-Pratt on P and text $T = \text{abacaabadcabacabaabb}$

j	0	1	2	3	4	5
P_j	a	b	a	c	a	b
$F(j)$	0	0	1	0	1	2



... Knuth-Morris-Pratt Algorithm

29/65

Construction of the failure function is similar to the KMP algorithm itself:

failureFunction(P):

Input pattern P of length m
Output failure function for P

```

F[0]=0
i=1, j=0
while i<m do
  if P[i]=P[j] then // we have matched j+1 characters
    F[i]=j+1
    i=i+1, j=j+1
  else if j>0 then // use failure function to shift P
    j=F[j-1]
  else
    F[i]=0 // no match
    i=i+1
  end if
end while
return F

```

Analysis of failure function computation:

- At each iteration of the while-loop, either
 - i increases by one, or
 - the "shift amount" $i-j$ increases by at least one (observe that $F(j-1) < j$)
- Hence, there are no more than $2 \cdot m$ iterations of the while-loop

⇒ failure function can be computed in $O(m)$ time

Analysis of Knuth-Morris-Pratt algorithm:

- Failure function can be computed in $O(m)$ time
- At each iteration of the while-loop, either
 - i increases by one, or
 - the "shift amount" $i-j$ increases by at least one (observe that $F(j-1) < j$)
- Hence, there are no more than $2 \cdot n$ iterations of the while-loop

⇒ KMP's algorithm runs in *optimal time* $O(m+n)$

Boyer-Moore vs KMP

Boyer-Moore algorithm

- decides how far to jump ahead based on the mismatched character in the text
- works best on large alphabets and natural language texts (e.g. English)

Knuth-Morris-Pratt algorithm

- uses information embodied in the pattern to determine where the next match could begin
- works best on small alphabets (e.g. A, C, G, T)

Tries

Preprocessing Strings

Preprocessing the *pattern* speeds up pattern matching queries

- After preprocessing P , KMP algorithm performs pattern matching in time proportional to the text length

If the text is large, immutable and searched for often (e.g., works by Shakespeare)

- we can preprocess the *text* instead of the pattern

A *trie* ...


- is a compact data structure for representing a set of strings

- e.g. all the words in a text, a dictionary etc.
- supports pattern matching queries in time proportional to the pattern size

Note: Trie comes from *retrieval*, but is pronounced like "try" to distinguish it from "tree"

Tries

Tries are trees organised using parts of keys (rather than whole keys)



Each node in a trie ...

- contains one part of a key (typically one character)
- may have up to 26 children
- may be tagged as a "finishing" node
- but even "finishing" nodes may have children

Depth d of trie = length of longest key value

Cost of searching $O(d)$ (independent of n)

Possible trie representation:


```
#define ALPHABET_SIZE 26
```

```
typedef struct Node *Trie;
```

```
typedef struct Node {  
    bool finish;        // last char in key?  
    Item data;          // no Item if !finish  
    Trie child[ALPHABET_SIZE];  
} Node;
```

```
typedef char *Key;
```

Note: Can also use BST-like nodes for more space-efficient implementation of tries



Trie Operations

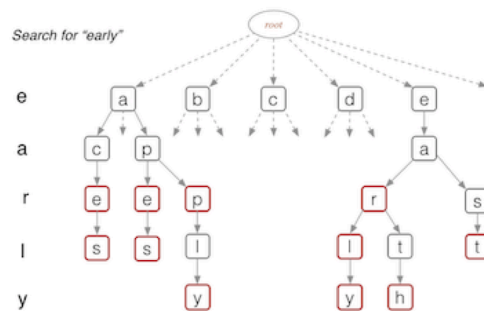
40/65

Basic operations on tries:

1. search for a key
2. insert a key

Trie Operations

41/65



... Trie Operations

42/65

Traversing a path, using char-by-char from Key:

```
find(trie, key):
    Input  trie, key
    Output pointer to element in trie if key found
           NULL otherwise

    node=trie
    for each char in key do
        if node.child[char] exists then
            node=node.child[char]    // move down one level
        else
            return NULL
        end if
    end for
    if node.finish then              // "finishing" node reached?
        return node
```

```
else
    return NULL
end if
```

... Trie Operations

43/65

Insertion into Trie:

```
insert(trie, item, key):
    Input  trie, item with key of length m
    Output trie with item inserted

    if trie is empty then
        t=new trie node
    end if
    if m=0 then
        t.finish=true, t.data=item
    else
        t.child[key[0]]=insert(trie, item, key[1..m-1])
    end if
    return t
```

... Trie Operations

44/65

Analysis of standard tries:

- $O(n)$ space
- insertion and search in time $O(d \cdot m)$
 - n ... total size of text (e.g. sum of lengths of all strings in a given dictionary)
 - m ... size of the string parameter of the operation (the "key")
 - d ... size of the underlying alphabet (e.g. 26)

Word Matching With Tries

Word Matching with Tries

46/65

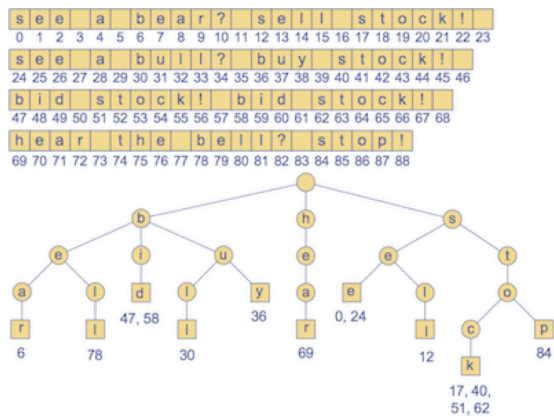
Preprocessing the text:

1. Insert all searchable words of a text into a trie
2. Each leaf stores the occurrence(s) of the associated word in the text

... Word Matching with Tries

47/65

Example text and corresponding trie of searchable words:




Compressed Tries

Compressed tries ...

- have internal nodes of degree ≥ 2
- are obtained from standard tries by compressing "redundant" chains of nodes

Example:




... Compressed Tries

Possible compact representation of a compressed trie to encode an array S of strings:

- nodes store *ranges of indices* instead of substrings
 - use triple (i, j, k) to represente substring $S[i][j..k]$
- requires $O(s)$ space ($s = \#$ strings in array S)


Example:



Pattern Matching With Suffix Tries


The *suffix trie* of a text T is the compressed trie of all the suffixes of T

Example:



... Pattern Matching With Suffix Tries

Compact representation:



... Pattern Matching With Suffix Tries

Input:

- compact suffix trie for text T
- pattern P

Goal:

- find starting index of a substrng of T equal to P

... Pattern Matching With Suffix Tries

`suffixTrieMatch(trie, P):`

Input compact suffix trie for text T , pattern P of length m
Output starting index of a substrng of T equal to P
 -1 if no such substrng exists

```
j=0, v=root of trie
repeat
  | // we have matched j+1 characters
```

```
if ∃w∈children(v) such that P[j]=T[start(w)] then
  i=start(w)           // start(w) is the start index of w
  x=end(w)-i+1         // end(w) is the end index of w
  if m≤x then          // length of suffix ≤ length of the node label?
    if P[j..j+m-1]=T[i..i+m-1] then
      return i-j       // match at i-j
    else
      return -1        // no match
  else if P[j..j+x-1]=T[i..i+x-1] then
    j=j+x, m=m-x       // update suffix start index and length
    v=w                // move down one level
  else return -1       // no match
end if
else
  return -1
end if
until v is leaf node
return -1              // no match
```

Analysis of pattern matching using suffix tries:

Suffix trie for a text of size n ...

- can be constructed in $O(n)$ time
- uses $O(n)$ space
- supports pattern matching queries in $O(s \cdot m)$ time
 - m ... length of the pattern
 - s ... size of the alphabet

Text Compression

Text Compression 56/65

Problem: Efficiently encode a given string X by a smaller string Y

Applications:

- Save memory and/or bandwidth

Huffman's algorithm

- computes frequency $f(c)$ for each character c
- encodes high-frequency characters with short code
- no code word is a prefix of another code word
- uses optimal *encoding tree* to determine the code words


Code ... mapping of each character to a binary code word

Prefix code ... binary code such that no code word is prefix of another code word

Encoding tree ...

- represents a prefix code
- each leaf stores a character
- code word given by the path from the root to the leaf (0 for left child, 1 for right child)

Example:




Text compression problem

Given a text T , find a prefix code that yields the shortest encoding of T

- short codewords for frequent characters
- long code words for rare characters

Example: $T = \text{abracadabra}$




T_1 requires 29 bits to encode text T ,

T_2 requires 24 bits

Huffman's algorithm

- computes frequency $f(c)$ for each character
- successively combines pairs of lowest-frequency characters to build encoding tree "bottom-up"

Example: abracadabra



Huffman Code

62/65

Huffman's algorithm using **priority queue**:


```
HuffmanCode(T):
  Input  string T of size n
  Output optimal encoding tree for T

  compute frequency array
  Q=new priority queue
  for all characters c do
    T=new single-node tree storing c
    join(Q,T) with frequency(c) as key
  end for
  while |Q| ≥ 2 do
    f1=Q.minKey(), T1=leave(Q)
    f2=Q.minKey(), T2=leave(Q)
    T=new tree node with subtrees T1 and T2
    join(Q,T) with f1+f2 as key
  end while
  return leave(Q)
```

... Huffman Code

63/65

Larger example: [a fast runner need never be afraid of the dark](#)



... Huffman Code

64/65

Analysis of Huffman's algorithm:

- $O(n+d \log d)$ time
 - n ... length of the input text T
 - s ... number of distinct characters in T

Summary

65/65

- Alphabets and words
- Pattern matching
 - Boyer-Moore, Knuth-Morris-Pratt
- Tries
- Text compression
 - Huffman code
- Suggested reading:
 - Tries ... Sedgewick, Ch.15.2

Produced: 17 Oct 2017